

Data Driven Resource Allocation for Distributed Machine Learning


Venkata Krishna Pillutla

www.cs.cmu.edu/~vpillutl

Thesis Committee

- Nina Balcan, Chair
- Alex Smola
- Christos Faloutsos

Collaborators




Machine Learning is Changing the World



“A breakthrough in machine learning would be worth ten
Microsofts”
(Bill Gates, Chairman, Microsoft)



“Machine learning is the next Internet”
(Tony Tether, former director, DARPA)




“Machine learning is the hot new thing”
(John Hennessy, President, Stanford)

The World is Changing ML



Outbreak of the “Data Epidemic”


New Applications



Introduction/Motivation

Machine Learning

- Traditional ML is centralized
- All the data is assumed to be on one machine



Distributed ML

Big Data in Google

-  **100 hours/min**
-  **100 petabytes**
-  **500+ million users**
-  **900+ million devices**



Distributed ML


Massive data is inherently distributed!



Also stored in a distributed manner. Eg: Yahoo! PNUTS


Distributed ML

In other cases, massive data centrally collected



Distributed ML

Communication: important resource (in addition to computation)



Typical Example: Learning Task

- Spam vs Not Spam

COMPOSE

Reading Group Presentation Inbox x

 **Travis Dick** <travis.barry.dick@gmail.com> Aug 9 

to learning-theor. ▾

Hi Everyone,

We postponed last weeks reading group meeting, so we will meet tomorrow at 3:00 in GHC 9115. Krishna will present an approximation algorithm for the capacitated k-center problem. The presentation will follow the paper "Centrality of Trees for Capacitated k-Center" available here: <http://arxiv.org/abs/1304.2983>.

Cheers,
Travis

 from [REDACTED] hide details 6:14 PM (10 minutes ago) Reply | ▾

<[REDACTED]@hotmail.com>
to [REDACTED]@yahoo.com
date Tue, May 5, 2009 at 6:14 PM
subject shopHello, friends
mailed-by hotmail.com

Retailer - Wholesaler Business
- PayPal.The Safer,Easier way to pay and get paid online.....

- Website:< http://www.leisureseries.com>


- The most fashionable branded shoes, clothes, jeans, handbags, wallets, belts, watches, jewelry, hat, sunglasses, ect.....Our business mainly brands such as: Gucci, Prada, Nike, LV.Puma.Chane,ACG, Balenciag,Nike,Chloe, Puma,Chanel,Dolce&Gabbana,MiuMiu,Mulberry,Fendi,and so on.....

- Thanks for your attention and welcome to enter my online store ,you will get more infomational.
- MSN: leisureseries@hotmail.com
- Email:leisureseries@yahoo.com.cn

- [leisureseries :\)](mailto:leisureseries :))


Another Example: Learning Task

Online advertising: Speed is of essence




Another Example: Learning Task

Online advertising: Speed is of essence




Another Example: Learning Task

Online advertising: Speed is of essence




Another Example: Learning Task


Online advertising: Speed is of essence



How to partition the data?



Random Partitioning






Random Partitioning


- Advantages
 - Easy to implement
 - Clean theory
- Disadvantages
 - Statistically sub-optimal
- Can we do better?




Our idea: Data dependent partitioning



Machine 1



Machine 2



Machine 3



Vapnik:
Locally Simple but
Globally Complex!



Pros and Cons

- Advantages
 - Distributed
 - More expressive concept class!
 - Better performance at same communication
- Possible Concern
 - More expressive dispatch rule is required




Data Dependent Partitioning

- How? Clustering




Data Dependent Partitioning

- For efficiency, cluster an initial sample




Requirements I

- Load-Balancing



Requirements II

- Fault-tolerance



Requirements III



- Efficient dispatch during deployment



Query

The Google logo, consisting of the word "Google" in its signature multi-colored font.

Users
(waiting for a real-time response)



Contributions*

- Balanced Clustering with Fault Tolerance
 - NP-hard
 - Approximation algorithm with strong guarantees
- Nearest Neighbor Dispatch
 - Efficient, Online Dispatch
 - Provably good
- Experiments
 - Classification accuracy after data dependent partitioning
 - Scalability

*Joint work with: Travis Dick, Mu Li, Colin White, Maria-Florina Balcan, Alex Smola
Under submission at AISTATS 2016



Balanced Clustering with Fault Tolerance

Requirements


<ul style="list-style-type: none">• Load balancing: Upper bound on cluster size: L fraction	Well studied [KS, ABC+, ABG+]
<ul style="list-style-type: none">• Load balancing: Lower bound on cluster size: l fraction	Not studied; very tricky
<ul style="list-style-type: none">• Fault tolerance: p replication	

Lower bounds are tricky


- Typically: OPT_k decreases as k increases
- With lower bounds:
 - Arbitrary number of local maxima [DLP+]



Handling Size Constraints



Handling Size Constraints



Algorithm Overview

- Notation:
 - y_i : point i is a center: opening
 - x_{ij} : point i is the center corresponding to j : assignments
 - V : set of points
- Works for any metric space (\mathcal{X}, d)

LP Relaxation

K -median: $c_{i,j} = d(i, j)$

y_i : opening

K -means: $c_{i,j} = d(i, j)^2$

x_{ij} : assignment

$$\min \sum_{i,j \in V} c_{ij} x_{ij}$$

$$\text{subject to: } \sum_{i \in V} x_{ij} = p, \quad \forall j \in V$$



$$\ell y_i \leq \sum_{j \in V} \frac{x_{ij}}{n} \leq L y_i, \quad \forall i \in V$$

$$\sum_{i \in V} y_i \leq k;$$

$$0 \leq x_{ij} \leq y_i \leq 1, \quad \forall i, j \in V.$$

LP Relaxation

- May open $2k$ half centers- requires rounding





Algorithm Overview

- Step 1 : Solve LP
- Step 2: Round opening
 - Greedy Coarse Clustering to get $\leq k$ coarse clusters: Monarch Procedure
 - Round centers locally within each coarse cluster
- Step 3: Round assignments
 - Round assignments globally with min-cost flow



Algorithm Overview: Step 1

- Solve LP
- Example: 8 points





Algorithm Overview: Step 2

- Perform coarse clustering: Monarch procedure
- Greedy
- Good guarantees





Algorithm Overview: Step 2

- Perform coarse clustering: Monarch procedure
- Greedy
- Good guarantees





Algorithm Overview: Step 2

- Round opening within each coarse cluster





Algorithm Overview: Step 2

- Round opening within each coarse cluster





Algorithm Overview: Step 2

- Round opening within cluster





Algorithm Overview: Step 3

- Round Assignments with min-cost flow.





Step 2: Monarch Procedure

- Greedily pick $\leq k$ points as monarchs





Step 2: Monarch Procedure

- Empires: Voronoi partitions about monarchs





Step 2: Monarch Procedure

- Greedy rule: pick point with highest contribution to the objective (as long as it does not have a monarch nearby)





Step 2: Monarch Procedure

- Why this greedy rule?




Step 2: Monarch Procedure Guarantees

- Points within an empire are close





Step 2: Monarch Procedure Guarantees

- Monarchs are far apart





Step 2: Monarch Procedure Guarantees

- Each empire has opening $\geq p/2$: Markov Inequality




Step 2: Monarch Procedure Guarantees

- Each empire has opening **at least 1!** (for $p > 1$)
- Round *locally* within each empire!




Step 2: Rounding within Empire

- Pick $\lfloor Y_{\mathcal{E}} \rfloor$ central points from each empire, each with opening $Y_{\mathcal{E}} / \lfloor Y_{\mathcal{E}} \rfloor$; make centers



$$Y_{\mathcal{E}} := \sum_{i \in \mathcal{E}} y_i$$



LP Relaxation

K -median: $c_{i,j} = d(i, j)$

y_i : opening

K -means: $c_{i,j} = d(i, j)^2$

x_{ij} : assignment

$$\min \sum_{i,j \in V} c_{ij} x_{ij}$$

$$\text{subject to: } \sum_{i \in V} x_{ij} = p, \quad \forall j \in V$$



$$\ell y_i \leq \sum_{j \in V} \frac{x_{ij}}{n} \leq L y_i, \quad \forall i \in V$$

$$\sum_{i \in V} y_i \leq k;$$

$$0 \leq x_{ij} \leq y_i \leq 1, \quad \forall i, j \in V.$$



Step 2: Rounding within Empire

- Same factor appears as violation of cluster size constraint: $Y_{\mathcal{E}}/\lfloor Y_{\mathcal{E}} \rfloor \leq p+2/p$




Step 2: Rounding Guarantee

- Obtain a feasible solution with integral y
- Cost bounded by triangle inequality




Step 3: Rounding assignments

- Easy: Min cost flow




Step 3: Rounding assignments

- Fractional LP solution implies a feasible flow




Step 3: Rounding assignments

- By Integral Flow Theorem, there is an optimal integral flow



Step 3: Rounding assignments

- Can be computed by standard algorithms



To sum up...


Theorem: There exist poly time approximation algorithms for balanced k -clustering with fault tolerance

- that output
 - 5 approx. for k -center
 - 11 approx. for k -median
 - 95 approx. for k -means, and
- cluster size constraint is violated by $(p+2)/p$
- replication between p and $p/2$.

Nearest Neighbor Dispatch


Requirements

- Dispatch a new point correctly and efficiently




Goal

- PAC Assumption: data are drawn iid from some fixed unknown distribution μ



Goal

- Given an iid sample from μ , cluster the distribution
- Balance constraints: Probability mass of each cluster is within (l, L) .



Our solution

- Cluster a sample (previous section)
- Extend clustering to the distribution
- How?

Clustering a Distribution


Assignments $f : \mathcal{X} \rightarrow \binom{k}{p}$

Centers $c : [k] \rightarrow \mathcal{X}$

K -median: $\min_{f,c} \quad \mathbb{E}_{x \sim \mu} \left[\sum_{i \in f(x)} \|x - c(i)\| \right]$


Find the Nearest Center?

- Doesn't work because of size constraints




An Idea: NN Extension

- Find nearest point from the original sample




An Idea: NN Extension

- Find nearest point from the original sample



An Idea: NN Extension

- Find nearest point from the original sample




NN Extension of a Clustering

Defined on distribution

$$\bar{g}_n(x) := g_n(NN_S(x))$$

Defined on sample



NN Extension

- Each point represents its Voronoi cell
- Sample level objective:

$$g_n : S \rightarrow \binom{k}{p}$$

$$c_n : [k] \rightarrow S$$

$$\min_{g_n, c_n} \sum_{j=1}^n w_j \left[\sum_{i \in g_n(x_j)} \|x_j - c_n(i)\| \right]$$

where $w_j = \mathbb{P}_{x \sim \mu}(NN_S(x) = x_j)$

NN Extension

- Weights are unknown
- Estimate weights from another sample drawn iid from μ .
- Cluster sample with estimated weights
- Use approx algo discussed earlier

NN Dispatch Algorithm

- Draw a second sample S' of size n' .
- Approximate weights w_j with estimates:

$$\hat{w}_j = \frac{|S' \cap V_j|}{n'}$$

- Find a balanced clustering (g_n, c_n) using estimated weights
- Return its NN extension

$$\bar{g}_n(x) = g_n(NN_S(x))$$

NN Dispatch

- Guarantee: NN Dispatch returns a good clustering of the distribution.
- Sub-optimality depends on
 - Quality of approximation on sample
 - Average ‘radius’ of Voronoi cell
$$\alpha(S) = \mathbb{E}_{x \sim \mu}(\|x - NN_S(x)\|)$$
 - Bias from returning clustering that are constant over Voronoi partitions
$$\beta(S) = \min_{h,c}(Q(\bar{h}, c) - Q(f^*, c^*))$$

s.t. h satisfies size constraints l, L


NN Dispatch

Theorem:

- If $n' = O((n + \ln 1/\delta)/\epsilon^2)$
- Algo on S returns solution within $r \cdot \mathcal{OPT} + s$
- Then w.p. $\geq 1 - \delta$
 - (\bar{g}_n, c_n) output satisfies sizes $(l - \epsilon, L + \epsilon)$
 - $Q(\bar{g}_n, c_n) \leq r \cdot Q(f^*, c^*) + s + 2(r + 1)pD\epsilon + p(r + 1) \cdot \alpha(S) + r \cdot \beta(S)$
 - $$\left(f^*, c^* = \mathcal{OPT}(l + \epsilon, L - \epsilon) \right)$$

NN Dispatch

- Can bound other terms
- Worst case exponential in dimension
 - Curse of dimensionality
- Better bounds with niceness assumptions
 - E.g., Doubling Measure



Experiments

Learning: Approximations

Balanced Clustering

- K-means++, with rebalancing

NN Dispatch

- Estimated weight = $1/n$
- Random Partition Trees for Approximate NN Search

Algorithm

- Cluster a small sample
- Extend the clustering to the rest of the training set with NN Dispatch
- Learn
 - independent model for each cluster or
 - in tandem, with partial or complete communication
- Testing
 - Query the appropriate model with NN Dispatch

Learning

- No communication:
 - Each cluster learns an independent model
 - Embarrassingly parallel
- Compare against:
 - Random partitioning with no communication
 - Random partitioning with full communication (global model)

Experimental Setup


- Run on a cluster with
 - 15 machines
 - 8 cores per machine, each of 2.4GHz
 - 32 GB shared memory per machine

Datasets



Datasets	Number of examples	Dimensionality
MNIST-8M	8 million	784
CIFAR-10-early	2.5 million	160
CIFAR-10-late	2.5 million	144
CTRc	0.8 million	232
CTRa	0.3 million	13 million
Criteo-Kaggle	45 million	34 million

CTR: Click Through Rate


Learning with no communication: MNIST-8M




Learning with no communication: CIFAR-10




Learning with no communication: CTRc



Learning with no communication: Criteo-Kaggle




Learning with no communication: Scalability




Learning with communication

- High dimensional datasets
- Feature occurrence: approx. power law



Learning with communication

- Tail features cannot be reliably learnt
- Scheme 1: Synchronize on tail features only across all clusters
- Scheme 2: Synchronize on all features, also store a local correction for head features
- Asynchronous Stochastic Gradient Descent



Scheme 1: Partial Communication

- Local model for head and synchronized model for the tail

$$f(\mathbf{x}) = \mathbf{w}_{i(\mathbf{x})} \cdot \mathbf{x}_h + \mathbf{w}_t \cdot \mathbf{x}_t$$

- Communication: not very high for relatively small size of head


Scheme 2: Full communication

- Each cluster stores a “correction” to the globally synchronized model


$$f(\mathbf{x}) = \mathbf{w}_{i(\mathbf{x})} \cdot \mathbf{x}_h + \mathbf{w}_g \cdot \mathbf{x}$$

- Communication: Equal to communication of fully synchronized global model

Performance on CTR data



How many local features?



Which scheme should I use?

- Dense data, images: No communication
- High dimensional data: With communication

Conclusion

- Data-dependent partitioning is good in both theory and practice!
- Balanced Clustering
- Nearest Neighbor Extension
- Experimental Evaluation

Thank You! Questions?

Collaborators

